



# Agent-based Systems

Paolo Turrini

🏠 [www.dcs.warwick.ac.uk/~pturrini](http://www.dcs.warwick.ac.uk/~pturrini) ✉ [p.turrini@warwick.ac.uk](mailto:p.turrini@warwick.ac.uk)

# The Plan

- **Logical Agents** [Week 1-2]
  - Knowledge, Preferences, Strategies and how to reason.
- **Decision Theory** [Week 3]
  - Probabilistic Beliefs and Expected Utility.
- **Game Theory** [Week 4-5]
  - Extensive Games and Opponent Modelling.
- **Learning Agents** [Week 6]
  - Markov Decision Processes, (Multi-Agent) Learning.
- **Collective Decision-Making** [Week 7-8]
  - Cooperation and Social Choice
- **Social Agents** [Week 9]
  - Coalitions, Matching, Social Networks.

# Markov Decision Processes

Policies are strategies

# Plan for Today

We now go back to a "typical" AI framework: Markov Decision Processes

- Plans and policies
- Optimal policies

These are "one player" games with perfect information.  
Except they are not played on trees.

This (and more) in RN's chapters 17-18.




**Stuart Russell and Peter Norvig**

Artificial Intelligence: a modern approach


2014 (3rd edition)

# The world



- Start at the starting square
- Move to adjacent squares
- Collision results in no movement
- The game ends when we reach either goal state  $+1$  or  $-1$


# The agent



The agent chooses between  $\{Up, Down, Left, Right\}$  and goes:

- to the intended direction with probability: e.g., 0.8
- to the left of the intended direction with probability: e.g., 0.1
- to the right of the intended direction with probability: e.g., 0.1

# Let's start walking





Walking is a repetition of throws:

- The probability that I walk right the first time:  $0.8$
- The probability that I walk right the second time:  $0.8$
- The probability that I walk right both times... is a product!  $0.8^2$

The environment is **Markovian**: the probability of reaching a state only depends on the state the agent is in and the action they perform.

It is also **fully observable**, like an extensive game (of imperfect information).

# Plans





$\{Up, Down, Left, Right\}$  denote the intended directions.

A **plan** is a finite sequence of **intended** moves, **from the start**.

So  $[Up, Down, Up, Right]$  is going to be the plan that, from the starting square, selects the intended moves in the specified order.



# Makings plans





**Goal:** get to  $+1$

Consider the plan  $[Up, Up, Right, Right, Right]$ .

- With deterministic agents, it gets us to  $+1$  with probability 1.
- But what happens to our stochastic agent instead?



What's the probability that  $[Up, Up, Right, Right, Right]$  gets us to  $+1$ ?

# Makings plans



- It's not  $0.8^5$ ! This is the probability that we get to  $+1$  when the plan works!

# Makings plans



- It's not  $0.8^5$ ! This is the probability that we get to  $+1$  when the plan works!
- The probability the plan does not work but still reaches  $+1$  is  $0.1^4 \times 0.8 = 0.00008$
- The correct answer is  $0.8^5 + 0.1^4 \times 0.8$
- Notice  $0.8^5 + 0.1^4 \times 0.8 < \frac{1}{3}$ , not great.

# Policies

$S^+$  is set of possible sequences of states (just like the histories of an extensive game!)

$A$  the set of available actions.

Then a **policy** is a function:

$$\pi : S^+ \rightarrow A$$

In words a policy is a protocol that at each possible decision point prescribes an action.

This **is** a strategy.

# A policy

3	→	→	→	+1
2	↑		↑	-1
1	↑	←	←	←
	1	2	3	4

This is a **state-based** policy. It recommends the same action at each state (so if two sequences end up with the same state, this policy is going to recommend the same action)

Now let's complicate things a little bit...

# Rewards



A **reward** function is a (utility) function of the form

$$r : S \rightarrow \mathbb{R}$$

All states, not just the terminal ones, get a reward!

Obviously, if you only care about terminal states, you may want to give zero to every other state. This is a more general model.

# Rewards





For instance, each non-terminal state:

- has 0 reward, i.e., only the terminal states matter;
- has negative reward, e.g., each move consumes  $-0.04$  of battery;
- has positive reward, e.g., I like wasting battery

Rewards are usually small, negative and uniform at non-terminal states. But the reward function allows for more generality.

# Rewards





Consider now the following. The reward is:  
 $+1$  at state  $+1$ ,  $-1$  at state  $-1$ ,  $-0.04$  in all other states.

What's the expected utility of  $[Up, Up, Right, Right, Right]$ ?



# Rewards





Consider now the following. The reward is:  
 $+1$  at state  $+1$ ,  $-1$  at state  $-1$ ,  $-0.04$  in all other states.

What's the expected utility of  $[Up, Up, Right, Right, Right]$ ?

IT DEPENDS

# Rewards



Consider now the following. The reward is:  
 $+1$  at state  $+1$ ,  $-1$  at state  $-1$ ,  $-0.04$  in all other states.

What's the expected utility of  $[Up, Up, Right, Right, Right]$ ?

IT DEPENDS on how we are going to put rewards together!

# Comparing states

Many ways of comparing states:

- summing all the rewards
- giving priority to the immediate rewards
- ...

# Utility of state sequences

There is only one general and 'reasonable' way to combine rewards over time.

**Discounted** utility function:  $u([s_0, s_1, s_2, \dots]) = r(s_0) + \gamma r(s_1) + \gamma^2 r(s_2) + \dots$

where  $\gamma \in [0, 1]$  is the **discounting factor**

Notice: **additive** utility function  $u([s_0, s_1, s_2, \dots]) = r(s_0) + r(s_1) + r(s_2) + \dots$  is just a discounted utility function where  $\gamma = 1$ .

# Discounting factor

$\gamma$  is a measure of the agent patience. How much more they value a gain of five today than a gain of five tomorrow, the day after etc...

- Used everywhere in AI, game theory, cognitive psychology
- A lot of experimental research on it
- Variants: discounting the discounting! I care more about the difference between today and tomorrow than the difference between some distant moment and the day after that!

# Discounting

- $\gamma = 1$  today is just another day
- $\gamma = 0$  today is all that matters

Basically  $\gamma$  is my attitude to risk towards the future!

Notice that stochastic actions introduce further gambling into the picture

# A problem

Here is a  $3 \times 101$  world.

50	-1	-1	-1	...	-1	-1	-1	-1
<i>s</i>				...				
-50	1	1	1	...	1	1	1	1

- start at *s*.
- two deterministic actions at *s*: either *Up* or *Down*
- beyond *s* you can only go *Right*.
- the numbers are the rewards you are going to get.

# A problem

Here is a  $3 \times 101$  world.

50	-1	-1	-1	...	-1	-1	-1	-1
<i>s</i>				...				
-50	1	1	1	...	1	1	1	1

- start at *s*.
- two deterministic actions at *s*: either *Up* or *Down*
- beyond *s* you can only go *Right*.
- the numbers are the rewards you are going to get.

Compute the expected utility of each action as a function of  $\gamma$



# Solution

The utility of  $Up$  is

$$50\gamma - \sum_{t=2}^{101} \gamma^t = 50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

# Solution

The utility of *Up* is

$$50\gamma - \sum_{t=2}^{101} \gamma^t = 50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

The utility of *Down* is

$$-50\gamma + \sum_{t=2}^{101} \gamma^t = -50\gamma + \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

# Solution

The indifference point is

$$50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma} = -50\gamma + \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

# Solution

The indifference point is

$$50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma} = -50\gamma + \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

Solving numerically, we have  $\gamma \approx 0.9844$ .

# Solution

The indifference point is

$$50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma} = -50\gamma + \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

Solving numerically, we have  $\gamma \approx 0.9844$ .

- If  $\gamma$  is strictly larger than this then *Down* is better than *Up*;

# Solution

The indifference point is

$$50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma} = -50\gamma + \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

Solving numerically, we have  $\gamma \approx 0.9844$ .

- If  $\gamma$  is strictly larger than this then *Down* is better than *Up*;
- If  $\gamma$  is strictly smaller than this then *Up* is better than *Down*;

# Solution

The indifference point is

$$50\gamma - \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma} = -50\gamma + \gamma^2 \frac{1 - \gamma^{100}}{1 - \gamma}$$

Solving numerically, we have  $\gamma \approx 0.9844$ .

- If  $\gamma$  is strictly larger than this then *Down* is better than *Up*;
- If  $\gamma$  is strictly smaller than this then *Up* is better than *Down*;
- Else, it does not matter.

# Markov Decision Process

A **Markov Decision Process** is a sequential decision problem for a:



# Markov Decision Process

A **Markov Decision Process** is a sequential decision problem for a:

- fully observable environment

# Markov Decision Process

A **Markov Decision Process** is a sequential decision problem for a:

- fully observable environment
- with stochastic actions

# Markov Decision Process

A **Markov Decision Process** is a sequential decision problem for a:



- fully observable environment
- with stochastic actions
- with a Markovian transition model

# Markov Decision Process

A **Markov Decision Process** is a sequential decision problem for a:

- fully observable environment
- with stochastic actions
- with a Markovian transition model
- and with discounted (possibly additive) rewards

# MDPs formally



## Definition

Let  $s$  be a state and  $a$  and action

Model  $P(s'|s, a)$  = probability that  $a$  in  $s$  leads to  $s'$

Reward function  $r(s)$  (or  $r(s, a)$ ,  $r(s, a, s')$ ) =  
$$\begin{cases} -0.04 & \text{(small penalty) for nonterminal states} \\ \pm 1 & \text{for terminal states} \end{cases}$$

# Expected utility of a policy


The expected utility (or value) of policy  $\pi$ , from state  $s$  is:

$$v^{\pi}(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r(S_t)\right]$$

$E$  is the expected utility of the sequences induced by:

- the policy  $\pi$  (the actions we are actually going to make)
- the initial state  $s$  (where we start)
- the transition model (where we can get to)

# Loops



- In principle we can go on forever!
- We are going to assume we need to keep going unless we hit a terminal state (**infinite horizon assumption**)

# Discounting

With discounting the utility of an infinite sequence is in fact **finite**.

If  $\gamma < 1$  and rewards are bounded above by  $\mathbf{r}$ , we have:

$$u[s_1, s_2, \dots] = \sum_{t=0}^{\infty} \gamma^t r(s_t) \leq \sum_{t=0}^{\infty} \gamma^t \mathbf{r} = \frac{\mathbf{r}}{1 - \gamma}$$



# Expected utility of a policy

An **optimal** policy (from a state) is the policy with the highest expected utility, starting from that state.

$$\pi_s^* = \operatorname{argmax}_{\pi} v^{\pi}(s)$$

We want to find the **optimal** policy.

# A remarkable fact

## Theorem

*With discounted rewards and infinite horizon*

$$\pi_s^* = \pi_{s'}^*, \text{ for each } s' \in S$$

*This means that the optimal policy does not depend on the sequences of states, but on the states only.*

*In other words, the optimal policy is a state-based policy.*

# A remarkable fact

## Theorem

*With discounted rewards and infinite horizon*

$$\pi_s^* = \pi_{s'}^*, \text{ for each } s' \in S$$

*This means that the optimal policy does not depend on the sequences of states, but on the states only.*

*In other words, the optimal policy is a state-based policy.*

**Idea:** Take  $\pi_a^*$  and  $\pi_b^*$ . If they both reach a state  $c$ , because they are both optimal, there is no reason why they should disagree (modulo indifference!). So  $\pi_c^*$  is identical for both (modulo indifference!). But then they behave the same at all states!

# Value of states

The **value of a state** is the value of the optimal policy from that state.

But then (VERY IMPORTANT): Given the values of the states, choosing the best action is just maximisation of expected utility!

**maximise the expected utility of the immediate successors**

# Value of states

3	0.812	0.868	0.912	<div>+1</div>
2	0.762		0.660	<div>-1</div>
1	0.705	0.655	0.611	0.388
	1	2	3	4

Figure: The values with  $\gamma = 1$  and  $r(s) = -0.04$

# Value of states

3	0.812	0.868	0.912	<div>+1</div>
2	0.762		0.660	<div>-1</div>
1	0.705	0.655	0.611	0.388
	1	2	3	4

Figure: The optimal policy

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) v(s')$$

Maximise the expected utility of the subsequent state

# Value of states




Figure: The optimal policy

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) v(s')$$

Maximise the expected utility of the subsequent state

# The Bellman equation

The definition of values of states, i.e., the expected utility of the optimal policy from there, leads to a simple relationship among values of neighbouring states:



# The Bellman equation

The definition of values of states, i.e., the expected utility of the optimal policy from there, leads to a simple relationship among values of neighbouring states:

**expected sum of rewards** =  
**current reward** +  $\gamma \times$  **expected sum of rewards after taking best action**

# The Bellman equation

Bellman equation (1957):

$$v(s) = r(s) + \gamma \max_a \sum_{s'} P(s' | (s, a)) v(s')$$

# The Bellman equation

Bellman equation (1957):

$$v(s) = r(s) + \gamma \max_a \sum_{s'} P(s' | (s, a)) v(s')$$

We can use it to compute the optimal policy!

# Value Iteration Algorithm

- 1 Start with arbitrary values
- 2 Repeat for every  $s$  simultaneously until “no change”

$$v(s) \leftarrow r(s) + \gamma \max_a \sum_{s'} v(s') P(s' \mid (s, a))$$

# Value Iteration Algorithm

- **Input**  $S$ ,  $A$ ,  $\gamma$ ,  $r$ , and  $P(s' \mid (s, a))$  for each  $s, s' \in S$ .
- **Input**  $\epsilon > 0$ , the error you want to allow
- **Output**  $v$ , the value of each state

# Value Iteration Algorithm

- 1 **Initialise**  $\delta_s := \epsilon \frac{(1-\gamma)}{\gamma}$  for all  $s$ ,  $v := 0$ , storing information to be updated

# Value Iteration Algorithm

- 1 **Initialise**  $\delta_s := \epsilon \frac{(1-\gamma)}{\gamma}$  for all  $s$ ,  $v := 0$ , storing information to be updated
- 2 **while**  $\delta_{s'} \geq \epsilon \frac{(1-\gamma)}{\gamma}$ , for some  $s'$ , **do**

# Value Iteration Algorithm

- 1 **Initialise**  $\delta_s := \epsilon \frac{(1-\gamma)}{\gamma}$  for all  $s$ ,  $v := 0$ , storing information to be updated
- 2 **while**  $\delta_{s'} \geq \epsilon \frac{(1-\gamma)}{\gamma}$ , for some  $s'$ , **do**
  - $v'(s) := r(s) + \gamma \max_a \sum_{s'} P(s' | (s, a))v(s')$



# Value Iteration Algorithm

- 1 **Initialise**  $\delta_s := \epsilon \frac{(1-\gamma)}{\gamma}$  for all  $s$ ,  $v := 0$ , storing information to be updated
- 2 **while**  $\delta_{s'} \geq \epsilon \frac{(1-\gamma)}{\gamma}$ , for some  $s'$ , **do**
  - $v'(s) := r(s) + \gamma \max_a \sum_{s'} P(s' | (s, a)) v(s')$
  - $\delta_s := |v'(s) - v(s)|$

# Value Iteration Algorithm

- ➊ **Initialise**  $\delta_s := \epsilon \frac{(1-\gamma)}{\gamma}$  for all  $s$ ,  $v := 0$ , storing information to be updated
- ➋ **while**  $\delta_{s'} \geq \epsilon \frac{(1-\gamma)}{\gamma}$ , for some  $s'$ , **do**
  - $v'(s) := r(s) + \gamma \max_a \sum_{s'} P(s' | (s, a)) v(s')$
  - $\delta_s := |v'(s) - v(s)|$
  - $v(s) := v'(s)$

# Value Iteration Algorithm

- ➊ **Initialise**  $\delta_s := \epsilon \frac{(1-\gamma)}{\gamma}$  for all  $s$ ,  $v := 0$ , storing information to be updated
- ➋ **while**  $\delta_{s'} \geq \epsilon \frac{(1-\gamma)}{\gamma}$ , for some  $s'$ , **do**
  - $v'(s) := r(s) + \gamma \max_a \sum_{s'} P(s' | (s, a)) v(s')$
  - $\delta_s := |v'(s) - v(s)|$
  - $v(s) := v'(s)$
- ➌ **Return**  $v$


# A fundamental fact

## Theorem


VIA:

- *terminates*
- *returns the optimal policy (for the input values)*

# VIA in action



# VIA in action



# VIA in action

3	0	0	0	<div>+1</div>
2	0		0	<div>-1</div>
1	0	0	0	0
	1	2	3	4

Initialise the values, for  $\gamma = 1, r = -0.04$

# VIA in action

3	0	0	0	<span style="border: 1px solid black; padding: 2px;">+1</span>
2	0		0	<span style="border: 1px solid black; padding: 2px;">-1</span>
1	0	0	0	0
	1	2	3	4


3	0	0	0	<span style="border: 1px solid black; padding: 2px;">+1</span>
2	0		0	<span style="border: 1px solid black; padding: 2px;">-1</span>
1	0	0	0	0
	1	2	3	4

Simultaneously apply the Bellmann update to all states

$$v(s) = r(s) + \gamma \max_a \sum_{s'} P(s' | (s, a)) v(s')$$




# Value Iteration Algorithm




# The state values

3	0.812	0.868	0.912	<div>+1</div>
2	0.762		0.660	<div>-1</div>
1	0.705	0.655	0.611	0.388
	1	2	3	4

# The optimal policy



# VIA in action




# VIA in action

3	0	0	0	<div>+1</div>
2	0		0	<div>-1</div>
1	0	0	0	0
	1	2	3	4

Initialise the values, for  $\gamma = 1, r = -0.4$

# VIA in action




# VIA in action

3	0	0	0	<div>+1</div>
2	0		0	<div>-1</div>
1	0	0	0	0
	1	2	3	4

Initialise the values, for  $\gamma = 1, r = -4$

# VIA in action






# VIA in action


3	0	0	0	<div>+1</div>
2	0		0	<div>-1</div>
1	0	0	0	0
	1	2	3	4

Initialise the values, for  $\gamma = 1, r = 0$


# VIA in action




# VIA in action




$r = [-\infty : -1.6284]$



$r = [-0.4278 : -0.0850]$



$r = [-0.0480 : -0.0274]$



$r = [-0.0218 : 0.0000]$

# Stocktaking

- Stochastic actions can lead to unpredictable outcomes
- But we can still find optimal "strategies", exploiting what happens in case we deviate from the original plan
- If we know what game we are playing and we play long enough...

**What next?** Learning in MDPs